

# Image Style Transfer using Convolutional Neural Networks

Clément Chadebec

ENS - MVA

January 20, 2020

## 1 Presentation of the method

## 2 Experiments

- Trade-off Content-Style
- Choice of layer in CNN
- Initialisation of Gradient Descent
- Stochastic results

# The method

Goal: Keep the content of an image  $\vec{p}$  while applying the style of an image  $\vec{a}$  to it



Figure: Content image  $\vec{p}$



Figure: Style image  $\vec{a}$

# Architecture of the Method

- Use of a pre-trained Convolutional Neural Network VGG19

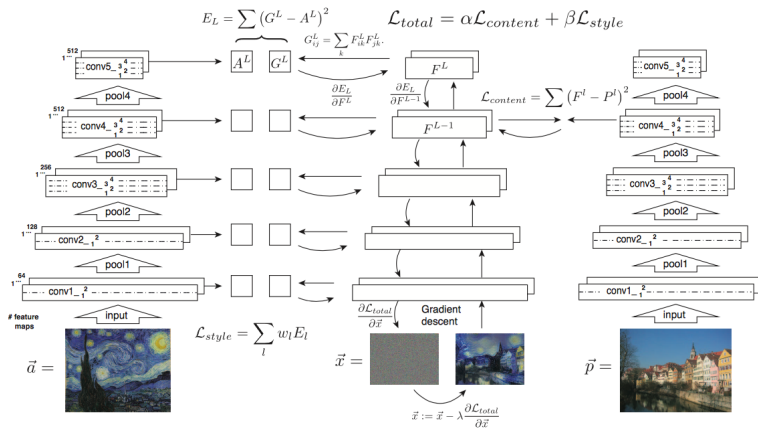


Figure: Architecture of the method

⇒ need for a representation of the images

# Content Representation

- Let  $P^l$  be the feature map of the content image  $\vec{p}$  produced by VGG19 at the output of layer  $l$
- Let  $F^l$  be the feature map of the image  $\vec{x}$  produced by VGG19 at the output of layer  $l$

$$\mathcal{L}_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{i,j}^l - P_{i,j}^l)^2$$

- The computation of the gradient gives

$$\frac{\partial \mathcal{L}_{content}}{\partial F_{i,j}^l} = \begin{cases} (F^l - P^l)_{i,j} & \text{if } F_{i,j}^l > 0 \\ 0 & \text{if } F_{i,j}^l < 0 \end{cases}$$

where  $F_{i,j}^l$  denotes the coefficient of the  $i$ -th filter (or feature) of the output of layer  $l$  at position  $j$

# Style Representation

- Several layers are considered
- Let  $E_l$  be the contribution of layer  $l$  in the loss
- Let  $w_l$  be a weight associated to the loss  $E_l$  in the total style loss
- We introduce the Gram matrix to compute  $E_l$

$$G_{i,j}^l = \sum_k F_{i,k}^l F_{j,k}^l$$

and we obtain

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{i,j}^l - A_{i,j}^l)^2$$

and finally we have the gradient

$$\frac{\partial \mathcal{L}_{style}}{\partial F_{i,j}^l} = \begin{cases} w_l \frac{1}{N_l^2 M_l^2} ((F^l)^\top (G^l - A^l))_{j,i} & \text{if } F_{i,j}^l > 0 \\ 0 & \text{if } F_{i,j}^l < 0 \end{cases}$$

# Trade-off Content - Style

- The total loss is as follows

$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

- Starting point  $\rightarrow$  White Gaussian noise
- number of iterations  $n = 3000$




Figure:  $\alpha/\beta = 10^{-1}$




Figure:  $\alpha/\beta = 10^{-2}$

# Trade-off Content - Style

- The total loss is as follows

$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

- Starting point  $\rightarrow$  White Gaussian noise
- number of iterations  $n = 3000$




Figure:  $\alpha/\beta = 10^{-3}$




Figure:  $\alpha/\beta = 10^{-4}$



# Layers' influence - Content

- Starting point  $\rightarrow$  White Gaussian noise
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: conv2\_2




Figure: conv4\_2

# Layers' influence - Style

- Starting point  $\rightarrow$  White Gaussian noise
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: conv2\_1



Figure: conv4\_1

# Layers' influence - Style

- Example of features map




Figure: "low" level layer




Figure: "high" level layer

# Initialisation

- Starting point  $\rightarrow$  White Gaussian noise
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Gaussian noise



Figure: Outcome

# Initialisation

- Starting point  $\rightarrow$  Content image
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$



Figure: Content image



Figure: Outcome

# Initialisation

- Starting point  $\longrightarrow$  Pathological case
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Pathological image




Figure: Outcome

# Initialisation

- Starting point  $\rightarrow$  Pathological case
- number of iterations  $n = 10000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Content image



Figure: Outcome

# Stochastic results

- Starting point  $\rightarrow$  Gaussian noise
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$



Figure: Outcome 1



Figure: Outcome 2



# Stochastic results

- Starting point  $\rightarrow$  Gaussian noise
- number of iterations  $n = 3000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Outcome 3




Figure: Outcome 4

# Photo realistic

- Starting point  $\rightarrow$  Content image
- number of iterations  $n = 5000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Dark image



Figure: Reference image

# Photo realistic

- Starting point  $\rightarrow$  Content image
- number of iterations  $n = 5000$
- weight factor  $\alpha/\beta = 10^{-3}$




Figure: Outcome

To be shown